

Data structures and Algorithms Searching

Pham Quang Dung

Hanoi, 2012

Outline

- 1 Sequential searches
- 2 Binary searches
- 3 Binary search trees
- 4 AVL trees
- 5 String searching
- 6 Map and Hashing

Sequential searches

- Input: a list $A[1..n]$ and an item x
- Output: position i such that $A[i] = x$

```
1 int sSearch(int A[], int L, int R, int x){  
  for(int i = L; i <= R; i++){  
    if(A[i]==x)  
      return i;  
  }  
  return -1;  
}
```

Outline

- 1 Sequential searches
- 2 Binary searches
- 3 Binary search trees
- 4 AVL trees
- 5 String searching
- 6 Map and Hashing

Binary search

- Input: a list $A[1..n]$ in non-decreasing order and an item x
- Output: position i such that $A[i] = x$

```
int bSearch(int A[], int L, int R, int x){  
2   int l = L;  
   int r = R;  
4   while(l <= r){  
       int mid = (l+r)/2;  
6       if(A[mid] == x)  
           return mid;  
8       if(A[mid] > x)  
           r = mid-1;  
10      else  
          l = mid+1;  
12  }  
   return -1;  
14 }
```


Outline

- 1 Sequential searches
- 2 Binary searches
- 3 Binary search trees**
- 4 AVL trees
- 5 String searching
- 6 Map and Hashing

Binary search trees

Binary search trees (BST)

- Each node has unique key
- Keys of nodes of left subtree of a node v are less than the key of v
- Keys of nodes of right subtree of a node v are greater than or equal to the key of v



Binary search trees

```
1 struct Node{  
2     int key;  
3     Node* leftChild;  
4     Node* rightChild;  
5 };  
6  
7 Node* root;
```


- `makeNode(int v)`: create a node with `key = v`, return the new created node
- `insert(Node* ptr, int v)`: insert a new node with `key = v` to the BST pointed by `ptr`
- `search(Node* ptr, int v)`: search a node having `key = v`, return this node if found
- `findMin(Node* ptr)`: return the node having smallest key of the BST pointed by `ptr`
- `del(Node* ptr, int v)`: remove a node having `key = v`

```
1 Node* makeNode(int x){  
    Node* p = new Node;  
3    p->key = x;  
    p->leftChild = NULL;  
5    p->rightChild = NULL;  
    return p;  
7 }
```

```
Node* insertNode(Node* r, int x){
2   if(r == NULL){
    r = makeNode(x);
4   } else if(r->key > x){
    r->leftChild = insertNode(r->leftChild, x);
6   } else if(r->key <= x){
    r->rightChild = insertNode(r->rightChild, x);
8   }
   return r;
10 }
```

```

Node* search(Node* r, int x){
2   if(r != NULL){
        if(r->key == x)
4           return r;
        else if(r->key > x)
6           return search(r->leftChild ,x);
        else
8           return search(r->rightChild ,x);
    }
10  return NULL;
}

```

```
1 Node* findMin(Node* r){
   if(r == NULL) return NULL;
3   Node* lmin = findMin(r->leftChild);
   if(lmin == NULL) return r;
5   if(lmin->key < r->key)
       return lmin;
7   else
       return r;
9 }
```

BST


```
1 Node* del(Node* r, int x){
  Node* tmp;
3  if(r == NULL) printf("Not Found\n");
  else if(x < r->key) r->leftChild = del(r->leftChild, x);
5  else if(x > r->key) r->rightChild = del(r->rightChild, x);
  else{// x = r->key, remove r means that make the smallest
        node of the right tree of r the root of the new tree
7      if(r->leftChild != NULL && r->rightChild != NULL){
        tmp = findMin(r->rightChild);
        r->key = tmp->key;
        r->rightChild = del(r->rightChild, tmp->key);
11     }else{
        tmp = r;// keep this node for freeing memory
13         if(r->leftChild == NULL)      r = r->rightChild;
        else if(r->rightChild == NULL)  r = r->leftChild;
15         delete tmp;
        }
17     }
    return r;
19 }
```

Outline


- 1 Sequential searches
- 2 Binary searches
- 3 Binary search trees
- 4 AVL trees**
- 5 String searching
- 6 Map and Hashing

AVL trees

- AVL is a BST (G. M. Adelson-Velskii and E. M. Landis, 1962)
 - the height of the left child differs from the height of the right child by at most 1 (balance property)
 - left and right subtrees are both AVL
- Modification (insertion, deletion of nodes) on AVL must conserve the balance property
- Assumption: the keys of all nodes are different (e.g., cannot construct an AVL for 8, 8, 9)



a. BST (not AVL)




b. AVL

AVL tree - recovery the balance property

- A BST T whose left and right subtrees are AVL
- Perform rotation actions so that the resulting BST is an AVL
- After the insertion or deletion of a node on an AVL
 - balance property of the AVL may loss
 - the height of any subtree change at most 1
 - identify which subtrees losing balance property


Recovery balance property

Case 1 (perform right rotation at K2)






Recovery balance property

Case 2 (perform left rotation at K1)






Recovery balance property

Case 3 (perform left rotation at K1 and then right rotation at K3)



Recovery balance property

Case 4 (perform right rotation at K3 and then left rotation at K1)




Outline

- 1 Sequential searches
- 2 Binary searches
- 3 Binary search trees
- 4 AVL trees
- 5 String searching**
- 6 Map and Hashing

String searching

- String matching problem: find one or all occurrences of a pattern in a given text
- Applications
 - information retrieval
 - Text editors
 - computational biology (DNA sequences)
- Formal formulation
 - A text is an array $T[1..n]$ and a pattern is an array $P[1..m]$ ($m \neq n$)
 - $T[i], P[j] \in$ a finite alphabet Σ (e.g., $\Sigma = \{0, 1\}$ or $\Sigma = \{a, \dots, z\}$)
 - We say that pattern P **occurs with shift s** in T if $0 \leq s \leq n - m$ and $T[s + 1..s + m] = P[1..m]$



String searching algorithms

- Naive
- Boyer-Moore
- Rabin-Karp
- Knuth-Morris-Pratt (KMP)



Algorithm 1: NaiveSM(P, T)

```
1 foreach  $s = 0..n-m$  do  
2    $i \leftarrow 1$ ;  
3   while  $i \leq m$  and  $P[i] = T[i + s]$  do  
4      $i \leftarrow i + 1$ ;  
5   if  $i \geq m$  then  
6     Output( $s$ );
```

Boyer-Moore algorithm

- Left to right shift
- Right to left scan
- Use information gained by preprocessing P in order to skip as many alignment as possible
- Bad character shift rule
 - $last[c]$: the right-most occurrence of c in P
 - When mismatch: shift P right by $\max\{j - last[c], 1\}$ where j is the position of mismatch character of P

Boyer-Moore algorithm



Boyer-Moore algorithm

```
1 void computeLast(){
2     for(int c = 0; c < 256; c++){
3         last[c] = 0;
4     }
5     for(int i = m; i >= 1; i--){
6         if(last[P[i]] == 0)
7             last[P[i]] = i;
8     }
9 }
10 void BoyerMoore(){
11     int s = 0;
12     while(s <= n-m){
13         int j = m;
14         while(j > 0 && T[j+s] == P[j])    j--;
15         if(j == 0){
16             Output(s);
17             s = s + 1;
18         } else {
19             int k = last[T[j+s]];
20             s = s + max(j-k, 1);
21         }
22     }
23 }
```

Rabin-Karp algorithm

- Convert the pattern $P[1..m]$ to a number:

$$p = P[1] * d^{m-1} + P[2] * d^{m-2} + \dots + P[m] * d^0$$

where each character $P[i]$ is viewed as a nonnegative integer $< d$, and d is the size of the alphabet

- Using Horner's rule:

$$p = P[m] + d * (P[m-1] + d * (\dots + d * P[1]) + \dots)$$

- Convert $T[s+1..s+m]$ to the integer

$$t_s = T[s+1] * d^{m-1} + \dots + T[s+m]$$

- **Note:** t_{s+1} can easily be computed from t_s as follows:

$$t_{s+1} = (t_s - T[s+1] * d^{m-1}) * d + T[s+m+1]$$

Rabin-Karp algorithm

- Drawback: when m is large, then the computation of p and t_s does not take constant time
- Solution: Compute p and t_s modulo a suitable number q
 - Still problem: $p \equiv t_s \pmod{q}$ does not mean that $p = t_s$, we have to check $P[1..m]$ and $T[s + 1..s + m]$ character by character to see if they are really identical
- Worst-case time is $\mathcal{O}(mn)$ where $P = a^m$ and $T = a^n$

Knuth-Morris-Pratt (KMP) algorithm

- Comparison: from left to right
- Shift: more than one position
- Preprocessing the pattern
 - Pattern $P[1..m]$
 - $\pi[q]$ is the length of the longest prefix of $P[1..m]$ which is also the **strictly** suffix of $P[1..q]$

Example

q	1	2	3	4	5	6	7	8	9	10
$P[q]$	a	b	a	b	a	b	a	b	c	a
$\pi[q]$	0	0	1	2	3	4	5	6	0	1


Knuth-Morris-Pratt (KMP) algorithm - preprocessing

```
1 void computePI(){
2     pi[1] = 0;
3     int k = 0;
4     for(int q = 2; q <= m; q++){
5         while(k > 0 && P[k+1] != P[q])
6             k = pi[k];
7         if(P[k+1] == P[q])
8             k = k + 1;
9         pi[q] = k;
10    }
11 }
```


Knuth-Morris-Pratt (KMP) algorithm - preprocessing

Denote $k = \pi[q]$


- If $P[q+1] = P[k+1]$, then $\pi[q+1] = \pi[q] + 1$



Knuth-Morris-Pratt (KMP) algorithm - preprocessing

Denote $k = \pi[q]$


- if $P[q+1] \neq P[k+1]$ and $P[q+1] = P[\pi[k]+1] = b$:
 - $P[1..k] = P[q-k+1..q] \Rightarrow P[k-\pi[k]+1..k] = P[q-\pi[k]+1..q]$
 - Moreover, $P[k-\pi[k]+1] = P[1..\pi[k]]$, so
 $P[1..\pi[k]] = P[q-\pi[k]+1..q]$,
 - Hence $P[1..\pi[k]+1] = P[q-\pi[k]+1..q+1]$, this means
 $\pi[q+1] = \pi[k]+1$



Knuth-Morris-Pratt (KMP) algorithm

```
1 void kmp(){
   int q = 0;
3  for(int i = 1; i <= n; i++){
      while(q > 0 && P[q+1] != T[i]){
5         q = pi[q];
      }
7      if(P[q+1] == T[i])
         q++;
9      if(q == m){
         cout << "match at position " << i-m+1 << endl;
11        q = pi[q];
      }
13  }
}
```

Knuth-Morris-Pratt (KMP) algorithm



Outline

- 1 Sequential searches
- 2 Binary searches
- 3 Binary search trees
- 4 AVL trees
- 5 String searching
- 6 Map and Hashing**

Map ADT

- Table stores key-value pairs
- Keys cannot be duplicated
- Operations
 - `size()`
 - `empty()`
 - `get(k)`: return the item having key k
 - `put(k, v)`: put a key-value pair (k, v) into the table
 - `remove(k)`: remove item having key k from the table
- Operations should be performed efficiently without sorting items of the table


Map implementation

- Arrays
- Linked-lists
- Binary search trees
- **Hash tables** (focus of this topic)

- A different approach to searching from the comparison-based methods
- Hashing tries to reference an item in the table directly based on its key without navigating through the table and comparing the search key with the keys of all items
- Hashing transform a key into a table address
- Two approaches
 - Direct-address tables
 - Hash tables

Direct-address tables

- Simple technique
- Work well when the universe U of keys is small
- Suppose each key is taken from $U = \{0, \dots, m - 1\}$ where m is not too large
- Use an array (**direct-address table**) $T[0..m - 1]$
 - Each slot k corresponds to a key k

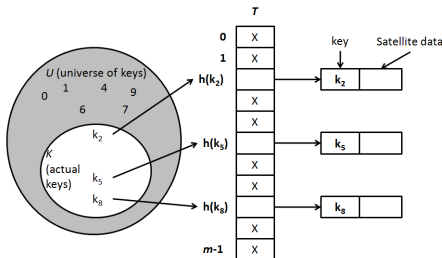


Hash table

- Each element with key k is stored in slot $h(k)$ of the hash table (called **hash function**)

$$h : U \rightarrow \{0, 1, \dots, m-1\}$$

- The size m of hash table is typically much less than the size of U
- Collision** when two keys k_1, k_2 hash the same slot: $h(k_1) = h(k_2)$
 - Solution by chaining
 - Solution by open addressing



Chaining

- Place all elements that hash to the same slot into the same linked list
- Slot i of the table stores a pointer to the head of the linked list

Algorithm 2: Put(k, v)

- 1 $x.key \leftarrow k$;
 - 2 $x.value \leftarrow v$;
 - 3 Insert x at the head of list $T[h(k)]$;
-

Algorithm 3: Get(k)

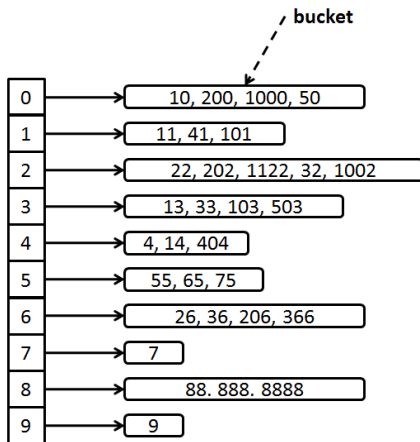
- 1 Search an element with key k in list $T[h(k)]$;
 - 2 **return** x ;
-

Algorithm 4: Remove(k)

- 1 $x \leftarrow \text{Get}(k)$;
 - 2 Delete x from list $T[h(k)]$;
-

Hashing: example

- keys: integer
- m is chosen to be 10
- $h(k) = k \bmod 10$
- Collision: two keys that hash to the same value (e.g., 22, 202)



Analysis of hashing with chaining

- We define a **load factor** $\alpha = \frac{n}{m}$
- Assumption of **simple uniform hashing**
- Average running time of unsuccessful search is $\Theta(1 + \alpha)$ (proof is detailed in “Introduction to Algorithms” book)
- Average running time of successful search is $\Theta(1 + \alpha)$ (proof is detailed in “Introduction to Algorithms” book)
- Put and get actions take $\mathcal{O}(1)$ if the lists are implemented as doubly linked lists

Open addressing

- All elements are stored in table itself
- When searching for an element: examine tables slots
- No list and no elements stored outside the table, avoid pointers together
- The hash table can fill up, no further insertion can be made

Open addressing - insertion

- Successively examine (**probe**) the table until an empty slot is found to put the key
- Instead of fixing the order $0, 1, m - 1$ ($\Theta(n)$ search time), the sequence of slots probed depends the key being inserted
- Extended hash function:
$$h : U \times \{0, 1, \dots, m - 1\} \rightarrow \{0, 1, \dots, m - 1\}$$
- The probe sequence for key k is $\langle h(k, 0), h(k, 1), \dots, h(k, m - 1) \rangle$ which is a permutation of $0, 1, \dots, m - 1$

Algorithm 5: Put(k, v)

```
1  $x.key \leftarrow k$ ;  
2  $x.value \leftarrow v$ ;  
3  $i \leftarrow 0$ ;  
4 while  $i < m$  do  
5    $j \leftarrow h(k, i)$ ;  
6   if  $T[j] = NULL$  then  
7      $T[j] \leftarrow x$ ;  
8     return  $j$ ;  
9    $i \leftarrow i + 1$ ;  
0 Error "hash table overflow";
```

Algorithm 6: $\text{Get}(k)$

```
1  $i \leftarrow 0$ ;  
2 while  $i < m$  do  
3    $j \leftarrow h(k, i)$ ;  
4   if  $T[j].\text{key} = k$  then  
5     return  $T[j]$ ;  
6    $i \leftarrow i + 1$ ;  
7 return NULL;
```

- **Uniform hashing** assumption: Each key is equally likely to have any of $m!$ permutation of $\{0, 1, \dots, m - 1\}$ as its probe sequence
- Three common techniques for probe sequence computation
 - Linear probing
 - Quadratic probing
 - Double hashing
- All of three techniques guarantee that $h(k, 0), h(k, 1), \dots, h(k, m - 1)$ is a permutation of $0, 1, \dots, m - 1$ for each key k
- None of three techniques fulfills the assumption of uniform hashing

Open addressing

- Linear probing: $h(k, i) = (h'(k) + i) \bmod m$ where h' is an ordinary hash function
- Quadratic probing: $h(k, i) = (h'(k) + c_1i + c_2i^2) \bmod m$ where h' is ordinary hash function and $c_1, c_2 \neq 0$ are constants
- Double hashing (one of the best method for open addressing):
 $(h(k, i) = (h_1(k) + ih_2(k)) \bmod m$ where h_1, h_2 are auxiliary hash functions
 - The value of $h_2(k)$ must be relatively prime to the hash table size m for the entire hash table to be searched (see “Introduction to Algorithms” for more detail)
 - Some approaches:
 - Let m be a power of 2 and design $h_2(k)$ such that it always returns an odd value
 - Let m be a prime and design $h_2(k)$ such that it always returns a positive integer less than m (e.g., $h_2(k) = 1 + k \bmod (m - 1)$)

Open addressing - analysis

- Inserting an element into an open-address hash table with load factor $\alpha = \frac{n}{m}$ requires at most $\frac{1}{1-\alpha}$ probes on average, assuming uniform hashing

Universal hashing

- If the chosen hash function is fixed, then there might be n keys that hash to the same slot, yielding an average retrieval time of $\Theta(n)$
- Solution: universal hashing
 - Select the hash function randomly from a carefully designed class of functions at the beginning of the execution
 - Randomization guarantees that no single input will always evoke worst-case behavior
 - Good average running time

Definition

Let \mathcal{H} be a finite collection of hash functions that map a given universe U of keys into the range $\{0, 1, \dots, m-1\}$. Such collection is said to be **universal** if for each pair of distinct keys $k, l \in U$, the number of hash functions $h \in \mathcal{H}$ for which $h(k) = h(l)$ is at most $\frac{|\mathcal{H}|}{m}$. In other words, for a given pair of distinct keys k, l , we pick a hash function from \mathcal{H} , the probability that $h(k) = h(l)$ is at most $\frac{1}{m}$.

Universal hashing

Theorem

Suppose that a hash function h is chosen randomly from a universal collection of hash functions and has been used to hash n keys into a table T of size m using chaining to resolve collisions. If key k is not in the table, then the expected length $E[n_{h(k)}]$ of the list that k hashes to is at most the load factor $\alpha = \frac{n}{m}$. If the key k is in the table, then the expected length $E[n_{h(k)}]$ of the list that k hashes to is at most $1 + \alpha$

Proof.

See “Introduction to Algorithms” book □

Universal hashing - example

- Choose a prime number p large enough so that every possible key k is in the range $0, \dots, p-1$, inclusive. Let $\mathbb{Z}_p = \{0, \dots, p-1\}$, and $\mathbb{Z}_p^* = \{1, 2, \dots, p-1\}$
- $p > m$ by assumption that the size of universe U is greater than the number of slots of the hash table
- Define hash function $h_{a,b}(k) = ((ak + b) \bmod p) \bmod m, \forall a \in \mathbb{Z}_p^*, b \in \mathbb{Z}_p$
- $\mathcal{H}_{p,m} = \{h_{a,b} : a \in \mathbb{Z}_p^*, b \in \mathbb{Z}_p\}$

Theorem

$\mathcal{H}_{p,m}$ defined above is **universal**

Proof.

See “Introduction to Algorithms” book □

Perfect hashing

- Hashing can provide excellent **worst-case** performance when the set of keys is **static**: once the keys are stored in the table, they never change
- **Perfect hashing**: $\mathcal{O}(1)$ memory access are required to perform a search in the worst-case
- Idea: Two-levels hashing with universal hashing at each level
 - First level: selected carefully from a family of universal hash functions
 - Second level uses hash tables instead of linked lists: Choose carefully hash function h_j for hash table S_j of slot j in order to guarantee that there are no collisions at the second level
 - Set the size m_j of hash table S_j to n_j^2 where n_j is the number of keys hashing to slot j

Perfect hashing

